

Task Description: Use Convolutional Neural Network (CNN) and Long-Short-Term Memory (LSTM) for sentiment classification.

Code: <https://github.com/salikadave/cse582-hw2>

Libraries used:

sklearn, pytorch, gensim

Dataset:

We are working with a subset of Yelp dataset that contains multiple reviews written by users from restaurants across multiple businesses.

Size of the dataset: 4.9GB

No of rows available for reviews: 6990280

Dataset Preprocessing:

```
review_final.head()
```

	review_id	business_id	review_stars	review_text
0	KU_O5udG6zpxOg-VcAEodg	XQfwVwDr-v0ZS3_CbbE5Xw	3.0	If you decide to eat here, just be aware it is...
1	BiTunyQ73aT9WBnpR9DZGw	7ATYjTlgM3jUlt4UM3IypQ	5.0	I've taken a lot of spin classes over the year...
2	saUsX_uimxRICVr67Z4Jig	YjUWPpI6HXG530lwP-fb2A	3.0	Family diner. Had the buffet. Eclectic assortm...
3	AqPFMleE6RsU23_auESxiA	kxX2SOes4o-D3ZQBkiMRfA	5.0	Wow! Yummy, different, delicious. Our favo...
4	Sx8TMOWLNuJBWer-0pcmoA	e4Vwtrqf-wpJfwesgvdxQ	4.0	Cute interior and owner (?) gave us tour of up...

Explored two methods of sentiment classification using data available about “stars” given by the user with the review text.

Method 1: (Used this in CNN)

Three classes: “negative” (stars < 2) ; “neutral” (stars == 3); “positive” (stars > 4);

Method 2: (Used this in LSTM)

Two classes: “negative” (stars <=2); “positive” (stars > 2);

For the purpose of this homework, the dataset of reviews was reduced to 10000 reviews for each of the classes.

Dataset split >> Training (75%) & Validation (25%)

Data Cleaning:

1. Removed newline characters
2. Remove punctuations and special characters
3. Convert all words to lowercase
4. Perform tokenization
5. Generate stemming tokens (used Porter Stemmer in CNN)

Embedding Vectors:

Used the Gensim word2vec embeddings for both CNN and LSTM.

Training Details:

CNN Training Time for 1 Epoch ~ 1m 8s

Embedding size = 500

of filters = 10

window sizes = 1,2,3,5

Activation function in CNN: tanh, relu, sigmoid

Loss function: Cross Entropy Loss

LSTM Training Time for 1 Epoch ~ 1m 20s

Due to time constraints experimentation is still in progress with the following:

clipping_quotient = 5 # set this value to avoid gradient explosion issues

num_epochs = 50

layers_count = 2

embedding_dimensions = 500

output_dimensions = 1

hidden_dimensions = 256

Activation function in LSTM: tanh, Sigmoid

Loss function: Binary Cross Entropy Loss

Overall, training time for CNN could be quicker than LSTM over 30 epochs.

Results:

CNN accuracy: 72%;

Modules	Parameters
convs.0.weight	5000
convs.0.bias	10
convs.1.weight	10000
convs.1.bias	10
convs.2.weight	15000
convs.2.bias	10
convs.3.weight	25000
convs.3.bias	10
fc.weight	120
fc.bias	3

	precision	recall	f1-score	support
0	0.77	0.76	0.77	2992
1	0.67	0.54	0.59	3044
2	0.71	0.87	0.78	2964
accuracy			0.72	9000
macro avg	0.72	0.72	0.71	9000
weighted avg	0.72	0.72	0.71	9000

LSTM accuracy: 69%

Accuracy comparison

CNN	tanh	72.66
	relu	72.85
LSTM	tanh	69.01

	relu	68.87
--	------	-------

Key Learning outcomes:

- CNN & LSTM implementation for sentiment analysis
- LSTM architecture implementation
- Usage of pytorch for building models
- Exploration of activation functions for binary and multi-class classification

Future Scope:

Experiment with more efficient pretrained models for word vector embeddings

Experiment with more layers with

References:

1. <https://towardsdatascience.com/sentiment-classification-using-cnn-in-pytorch-fba3c6840430>
2. Building LSTM classes for Hebrew text classification:
<https://galhever.medium.com/sentiment-analysis-with-pytorch-part-4-lstm-bilstm-model-84447f6c4525>
3. Understanding LSTM from scratch:
<https://medium.com/hackernoon/understanding-architecture-of-lstm-cell-from-scratch-with-code-8da40f0b71f4>